ANTI-MONEY LAUNDERING · 2026 EDITION

An AML program every decision can be replayed.

Process and governance gaps outnumber data and model gaps roughly 2:1 in every recent consent order. The cost of being unable to evidence a decision is now larger than the cost of missing one.

PREPARED FOR **CCO · MLRO · AUDIT COMMITTEE**

Source: FCA Dear-CEO Letter (Mar 2024) · FinCEN BSA consent order (Oct 2024 · $3.09B) · LexisNexis True Cost of Compliance (Feb 2024)

Process and governance gaps account for ~2:1 of every consent-order finding — not detection failures.

Five of the largest 2024–25 enforcement actions surveyed. Regulators rarely allege the bank missed a typology. They allege the bank cannot evidence what it did.

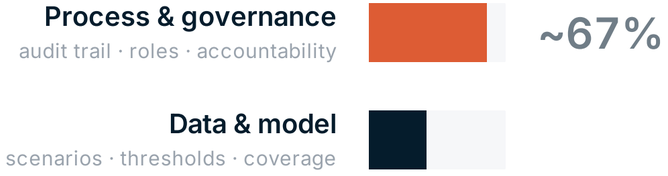

> *"Decisions made in relation to financial crime were not supported by evidence or an audit trail of debate and challenge."*
>
> **FCA DEAR-CEO LETTER TO ANNEX 1 FIRMS** · MARCH 2024 · STILL OPERATIVE IN 2026

KEY TAKEAWAY Lead with proof, not detection. The single highest-leverage place to spend a compliance dollar in 2026 is on *provability*.

Source: Synthesis of the five largest 2024–25 BSA / AML consent orders · FCA Dear-CEO Letter Mar 2024 · K&L Gates 2024 BSA enforcement-action review

The bill arrives in three sizes — and the largest is still growing.

A single record-breaking penalty in 2024, an industry-wide annual run-rate, and a daily-process tax — the same dollars that buy detection also buy poor evidence and poor signal.

SINGLE-BANK PENALTY · 2024

$3.09B

Largest BSA enforcement action of 2024. Includes a U.S. asset cap (~$434B) — only the second ever on a major bank.

FinCEN · DOJ · OCC consent order Oct 2024

ANNUAL INDUSTRY SPEND · US + CA

$61B

Annual cost of financial-crime compliance. The CCO's recurring board question — *"is the juice worth the squeeze?"* — is now FinCEN's literal RFI title.

LexisNexis Risk Solutions Feb 2024 · FinCEN Sep 2025 RFI

INDUSTRY FALSE-POSITIVE BASELINE

95%

Of every 100 alerts, 95 close as legitimate. Analyst attention is the scarce resource — and most of it is spent on noise, not signal.

Flagright industry analysis · NICE Actimize 2026

KEY TAKEAWAY The visible cost is the fine. The invisible cost is the analyst hour spent on the 95th false positive — every day, every quarter, every year.

Source: FinCEN consent-order release Oct 2024 · LexisNexis True Cost of FCC Feb 2024 · FinCEN AML-cost RFI Sep 2025 · Flagright TM industry analysis

The diagnosis is in the language: four sentences a CCO says before the regulator does.

Quoted verbatim from public regulator letters and consent orders — none paraphrased. Every line below is something a Chief Compliance Officer or regulator has put on the record between 2024 and 2026.

PAIN-1 · AUDIT-DEFENSIBILITY

"Decisions made in relation to financial crime were not supported by evidence or an audit trail of debate and challenge."

FCA · Dear-CEO Letter to Annex 1 firms · Mar 2024

PAIN-2 · BACKLOG VISIBILITY

"Examiner testing only determined whether controls existed rather than if they were actually being used."

K&L Gates · 2024 BSA enforcement-action review

PAIN-6 · MODEL VALIDATION

"Transaction monitoring systems... meet the definition of 'model' and are held to the same standard."

ACAMS / 2021 Interagency Statement · still cited 2026

PAIN-8 · ROI

"Is the juice worth the squeeze?"

FinCEN · AML-Compliance-Cost RFI literal subject line · Sep 2025

Source: All quotes verbatim from cited regulator publications · full pain-point research at *docs/research/2026-04-aml-process-pain.md*

One specification serves three reviewers — without re-translation.

One machine-readable file. Authored by Compliance, executed unchanged by Engineering, replayed byte-for-byte by an Auditor — every quarter.

SOURCE OF TRUTH

aml.yaml

written once · committed to version control

FIRST LINE · COMPLIANCE

The risk policy in machine-readable form, in our own language.

"Our scenarios, our thresholds, every line cited to a regulation."

ENGINEERING

The schema we build from. The data the engine consumes.

"Compiles deterministically — same spec, same data, same answer."

AUDITOR · 2LOD

The artifact I can replay any quarter, byte-for-byte.

"When the regulator asks 'who decided?' — the answer is in the file."

KEY TAKEAWAY Remove the translation layer between 1LoD / 2LoD / 3LoD — and you remove the most-cited finding of 2024–26.

Source: AML Open Framework specification model · primary-source quotes from FCA Dear-CEO 2024 + K&L Gates 2024 BSA review

Same workflow, two days. The unmeasured win is the analyst hours the framework gives back.

Detection accuracy gets the headline; the line item that drains the floor is what an analyst does *between* alerts.

TODAY · WITHOUT THE FRAMEWORK

ALERT TRIAGE

Open **8 systems** for one alert — TM, KYC, sanctions, CRM, transactions, prior cases — and copy-paste into a Word file. *~95% close as legitimate.*

KYC / PERIODIC REVIEW

Pull the customer dossier from 3 systems, attach screenshots, justify the risk rating in a free-text box. Annual review backlog grows by month-end.

STR DRAFTING

Hunt the customer history, the transaction trail, the typology citation. Re-type the narrative in goAML. The investigation is in the analyst's head.

BOI REFRESH

Email the relationship manager, chase a beneficial-ownership form, attach to the case file. Stale-BOI list is in someone's spreadsheet.

BACKLOG STATUS

"How many alerts are over SLA?" — *"I'll ask the team."* Dashboard shows last week's red. Board approves the program anyway.

WITH THIS FRAMEWORK · SAME TUESDAY

ALERT TRIAGE

One evidence bundle. Customer record, transactions, prior alerts, network neighbours, regulator citation — all chained, hash-verified, on one screen.

KYC / PERIODIC REVIEW

Risk-rating computed from the spec, dossier auto-assembled, EDD-due list rendered live from Customer Risk page. Reviews stop being a backlog.

STR DRAFTING

Narrative auto-drafted from the alert evidence chain. Analyst reviews and signs — doesn't re-type. Filing-latency p95 is a live metric.

BOI REFRESH

BOI Workflow page shows every reporting-company customer + freshness state. Stale BOI surfaces as an alert, not as someone's spreadsheet.

BACKLOG STATUS

SLA timer + escalation engine. Per-queue breach-rate badges. Red queues become loud, not slide-decoration. CCO sees real status first.

KEY TAKEAWAY The unmeasured win isn't accuracy — it's the *analyst hours* the team gets back. Multiply across 100 analysts, 50 weeks: same dollars, signal instead of motion.

Source: industry false-positive baseline · Flagright + NICE Actimize 2026 · STR-filing-latency p95 cited in 2024 BSA consent orders · framework workflows in *Investigations*, *BOI Workflow*, *Tuning Lab*, *Audit & Evidence* dashboard pages

The other half of the win is political. When the spec is the artifact, cross-team coordination collapses.

Most of an AML program's calendar is consumed by handoffs — 1LoD ⇄ 2LoD ⇄ IT ⇄ MRM ⇄ Audit. None of that produces detection. It produces alignment, and only when the alignment holds.

WHY THIS MATTERS

"Examiner testing only determined whether controls existed rather than if they were actually being used."

Across the largest 2024–25 enforcement actions, the same pattern recurs: **the controls existed; the coordination didn't.** Internal alerts get filed and never escalated. Rule changes wait on a calendar. Audit prep starts at the regulator's email. Removing the meetings removes the failure mode.

K&L GATES · 2024 BSA ENFORCEMENT-ACTION REVIEW · REFERENCED IN FCA DEAR-CEO MAR 2024

MONTHS → MINUTES

Rule-change approval → **one PR with cited regulation**

Today: months of 1LoD ⇄ 2LoD ⇄ IT ⇄ Audit alignment. With the framework: *git diff,* reviewer signs, audit chain records the change.

BRA refresh → **live render from spec**

Today: quarterly committee assembling spreadsheets. With the framework: the spec *is* the BRA — *Framework Alignment* dashboard shows it live.

Annual MRM challenge → **backtester self-serves**

Today: 4-week consultant engagement to validate scenarios. With the framework: per-rule precision/recall over 4 quarters, generated on demand.

Pre-exam evidence assembly → `aml audit-pack`

Today: 8-person sprint, 6 weeks, "what did we run?" forensics. With the framework: one CLI command exports the regulator-ready ZIP.

KEY TAKEAWAY Detection vendors sell better alerts. The framework removes the *meetings about the alerts* — and that is where the AML program's calendar actually lives.

Source: K&L Gates 2024 BSA enforcement-action review · FCA Dear-CEO Letter Mar 2024 · framework workflows in *Spec Editor*, *Framework Alignment*, *Tuning Lab*, *Audit Evidence* dashboard pages · `aml audit-pack` CLI

Each pain has a single, evidenced answer — one file away.

Direct mapping from public-record regulator findings to specific framework capabilities. No black box, no vendor add-ons.

	LEADER'S PAIN (REGULATOR-CITED)		FRAMEWORK CAPABILITY (WHAT CHANGES)
1	*"We can't prove what we did."*	→	**Hash-chained** decisions.jsonl + deterministic replay. Any historical run reproducible byte-for-byte.
2	*"The backlog has been red on the board for years."*	→	**SLA timer + escalation engine.** Red queues become loud, not slide-decoration. Per-queue breach-rate badges visible to the board.
3	*"Our risk assessment is a PDF nobody trusts."*	→	**The spec IS the risk assessment** — machine-readable, cited line-by-line. *Framework Alignment* dashboard renders the live BRA from the spec.
4	*"My monitoring system is a model and I cannot validate it."*	→	**Per-rule MRM dossier + 4-quarter backtester.** "Is rule X still earning its keep?" answered before lunch — no vendor study.

Source: Pain points 1, 2, 4, 6 from *docs/research/2026-04-aml-process-pain.md* · capability mapping in framework README + audit-evidence.md

The deterministic + spec-as-source-of-truth corner is uncontested.

Plotted against six 2026 buyer archetypes on the two axes regulators care about — *completeness of vision* and *ability to execute*.

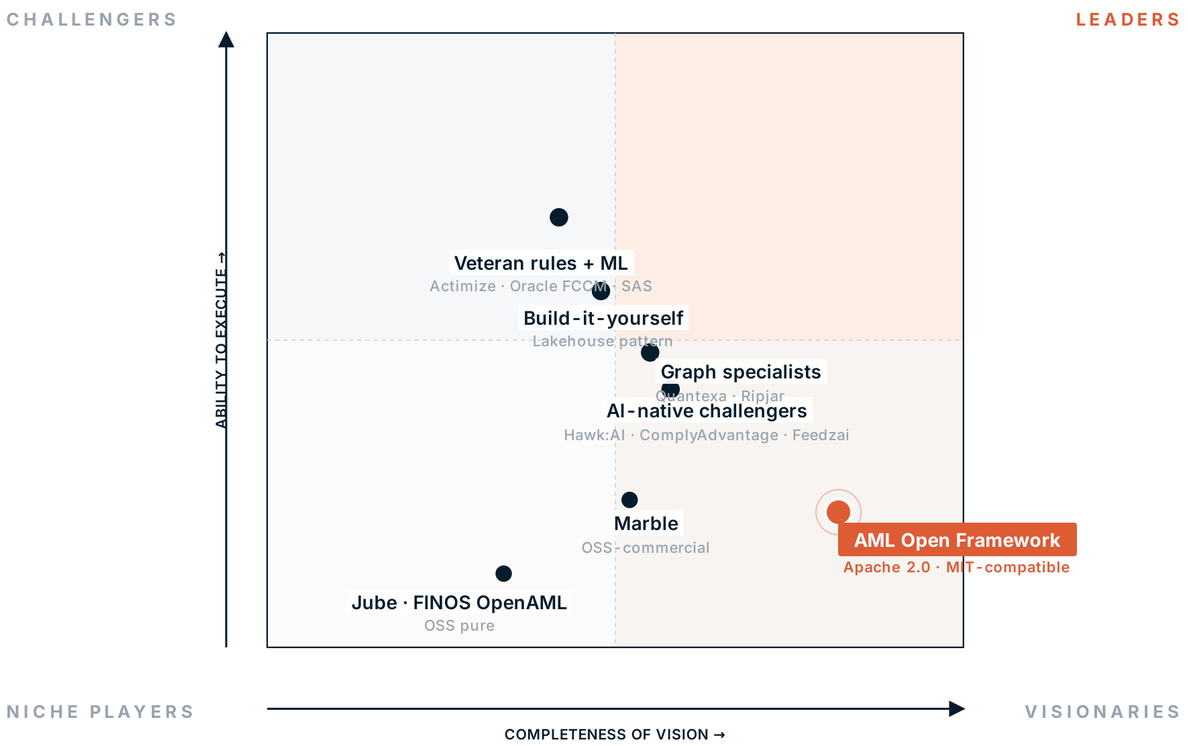

STRATEGIC PLANNING ASSUMPTION

By April 2026, SR 26-2 folds AML transaction monitoring into model-risk inventory — every **ML-first vendor carries a model-validation tax** on every typology.

- **AML Open Framework — Visionary, by design**
  Spec-as-policy is *not a model* under SR 26-2. Hash-chained replay; multi-jurisdiction templating; ISO 20022 in-tree. *Buyer that lands:* 2LoD MRM / Audit at any Cat-2+ bank.
- **Veteran platforms — Leaders by inertia**
  Win Tier-1 RFPs by default. Differentiator: scale + audit comfort. *Where we slot:* 2LoD challenger model for SR 26-2 independent re-implementation.
- **AI-native challengers — Visionary but ML-taxed**
  Hawk:AI Series C ($56M+$10M); ComplyAdvantage cites 65–85% FP reduction. *Adjacent, not opposed* — they win mid-market; we cover the same buyer only when the buyer wants source-available code.

KEY TAKEAWAY The framework doesn't compete with veteran platforms on detection scale — it competes on *provability*. Under SR 26-2, that is now the buyer's procurement criterion.

Source: Framework analysis based on cited 2026 vendor information (SymphonyAI, Alessa, Tookitaki, ComplyAdvantage, SiliconANGLE) · SR 26-2 letter Apr 2026 · OCC Bulletin 2026-13 · *Not a Gartner publication.*

The framework arrives built, not blueprinted.

Open-source · runs in your perimeter · no per-seat licence · every rule cites a regulation · re-runs deterministically.

8

CURATED TYPOLOGIES

Every line cited to a regulator publication. Day-one detection.

4

ACTIVE JURISDICTIONS

US (FinCEN) · CA (FINTRAC) · UK (FCA) · EU (EBA). FATF Travel Rule.

26

DASHBOARD PAGES

Persona-aware. Sidebar filters by role — CCO sees board signal, analyst sees queue.

9

REFERENCE SPECS

Schedule-I bank · community bank · VASP · cyber-fraud · trade-based ML · …

1,632

TESTS · ALL GREEN

5 CI jobs per pull request. Lint, unit, API, e2e dashboard, docker.

SHA-256

AUDIT CHAIN

Hash-chained decision log. Tampering breaks attestation, by construction.

MIT

LICENCE

Apache 2.0 / MIT compatible. Forkable. Runs entirely in your perimeter.

$0

PER-SEAT FEES

No vendor licence renewals. Effectiveness pack quantifies what you spent.

Source: Live counts from *main* branch as of 2026-04-29 · github.com/tomqwu/aml_open_framework

THREE PATHS FROM HERE

The next step takes five minutes.

READ · 15 MIN

Show me the research.

Open the 10-pain-point doc with primary regulator citations plus a 2024 enforcement case study. Frame is unbranded.

TRY · 5 MIN

Run it on my desk.

Clone, install, run a real audit pack against the canonical spec. No vendor demo cycle.

TALK · 30 MIN

Walk me through it.

Bring your spec or pain — pull-request a typology, open a Discussion, or tag the maintainer.

OPEN SOURCE · MIT

github.com/tomqwu/aml_open_framework

Source: AML Open Framework · primary-source research grounded in FCA / FinCEN / OSFI / FINTRAC public record