ANTI-MONEY LAUNDERING · 2026 EDITION

# An AML program every decision can be replayed.

Process and governance gaps outnumber data and model gaps roughly 2:1 in every recent consent order. The cost of being unable to evidence a decision is now larger than the cost of missing one.

PREPARED FOR **CCO · MLRO · AUDIT COMMITTEE**

**Source:** FCA Dear-CEO Letter (Mar 2024) · FinCEN BSA consent order (Oct 2024 · $3.09B) · LexisNexis True Cost of Compliance (Feb 2024)

# Process and governance gaps account for ~2:1 of every consent-order finding — not detection failures.

Five of the largest 2024–25 enforcement actions surveyed. Regulators rarely allege the bank missed a typology. They allege the bank cannot evidence what it did.

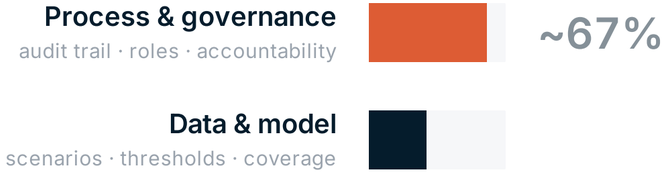


> *"Decisions made in relation to financial crime were not supported by evidence or an audit trail of debate and challenge."*
>
> **FCA DEAR-CEO LETTER TO ANNEX 1 FIRMS** · MARCH 2024 · STILL OPERATIVE IN 2026

**KEY TAKEAWAY** Lead with proof, not detection. The single highest-leverage place to spend a compliance dollar in 2026 is on *provability*.

**Source:** Synthesis of the five largest 2024–25 BSA / AML consent orders · FCA Dear-CEO Letter Mar 2024 · K&L Gates 2024 BSA enforcement-action review

# The bill arrives in three sizes — and the largest is still growing.

A single record-breaking penalty in 2024, an industry-wide annual run-rate, and a daily-process tax — the same dollars that buy detection also buy poor evidence and poor signal.

SINGLE-BANK PENALTY · 2024

## $3.09B

Largest BSA enforcement action of 2024. Includes a U.S. asset cap (~$434B) — only the second ever on a major bank.

FinCEN · DOJ · OCC consent order Oct 2024

ANNUAL INDUSTRY SPEND · US + CA

## $61B

Annual cost of financial-crime compliance. The CCO's recurring board question — *"is the juice worth the squeeze?"* — is now FinCEN's literal RFI title.

LexisNexis Risk Solutions Feb 2024 · FinCEN Sep 2025 RFI

INDUSTRY FALSE-POSITIVE BASELINE

## 95%

Of every 100 alerts, 95 close as legitimate. Analyst attention is the scarce resource — and most of it is spent on noise, not signal.

Flagright industry analysis · NICE Actimize 2026

**KEY TAKEAWAY** The visible cost is the fine. The invisible cost is the analyst hour spent on the 95th false positive — every day, every quarter, every year.

**Source:** FinCEN consent-order release Oct 2024 · LexisNexis True Cost of FCC Feb 2024 · FinCEN AML-cost RFI Sep 2025 · Flagright TM industry analysis

# The diagnosis is in the language: four sentences a CCO says before the regulator does.

Quoted verbatim from public regulator letters and consent orders — none paraphrased. Every line below is something a Chief Compliance Officer or regulator has put on the record between 2024 and 2026.

PAIN-1 · AUDIT-DEFENSIBILITY

"Decisions made in relation to financial crime were not supported by evidence or an audit trail of debate and challenge."

**FCA** · Dear-CEO Letter to Annex 1 firms · Mar 2024

PAIN-2 · BACKLOG VISIBILITY

"Examiner testing only determined whether controls existed rather than if they were actually being used."

**K&L Gates** · 2024 BSA enforcement-action review

PAIN-6 · MODEL VALIDATION

"Transaction monitoring systems... meet the definition of 'model' and are held to the same standard."

**ACAMS / 2021 Interagency Statement** · still cited 2026

PAIN-8 · ROI

"Is the juice worth the squeeze?"

**FinCEN** · AML-Compliance-Cost RFI literal subject line · Sep 2025

**Source:** All quotes verbatim from cited regulator publications · full pain-point research at *docs/research/2026-04-aml-process-pain.md*

# One specification serves three reviewers — without re-translation.

One machine-readable file. Authored by Compliance, executed unchanged by Engineering, replayed byte-for-byte by an Auditor — every quarter.

SOURCE OF TRUTH

**aml.yaml**

written once · committed to version control

### FIRST LINE · COMPLIANCE

The risk policy in machine-readable form, in our own language.

*"Our scenarios, our thresholds, every line cited to a regulation."*

### ENGINEERING

The schema we build from. The data the engine consumes.

*"Compiles deterministically — same spec, same data, same answer."*

### AUDITOR · 2LOD

The artifact I can replay any quarter, byte-for-byte.

*"When the regulator asks 'who decided?' — the answer is in the file."*

**KEY TAKEAWAY** Remove the translation layer between 1LoD / 2LoD / 3LoD — and you remove the most-cited finding of 2024–26.

**Source:** AML Open Framework specification model · primary-source quotes from FCA Dear-CEO 2024 + K&L Gates 2024 BSA review

# Same workflow, two days. The unmeasured win is the analyst hours the framework gives back.

Detection accuracy gets the headline; the line item that drains the floor is what an analyst does *between* alerts.

## TODAY · WITHOUT THE FRAMEWORK

### ALERT TRIAGE

Open **8 systems** for one alert — TM, KYC, sanctions, CRM, transactions, prior cases — and copy-paste into a Word file. *~95% close as legitimate.*

### KYC / PERIODIC REVIEW

Pull the customer dossier from 3 systems, attach screenshots, justify the risk rating in a free-text box. Annual review backlog grows by month-end.

### STR DRAFTING

Hunt the customer history, the transaction trail, the typology citation. Re-type the narrative in goAML. The investigation is in the analyst's head.

### BOI REFRESH

Email the relationship manager, chase a beneficial-ownership form, attach to the case file. Stale-BOI list is in someone's spreadsheet.

### BACKLOG STATUS

"How many alerts are over SLA?" — *"I'll ask the team."* Dashboard shows last week's red. Board approves the program anyway.

## WITH THIS FRAMEWORK · SAME TUESDAY

### ALERT TRIAGE

One evidence bundle. Customer record, transactions, prior alerts, network neighbours, regulator citation — all chained, hash-verified, on one screen.

### KYC / PERIODIC REVIEW

Risk-rating computed from the spec, dossier auto-assembled, EDD-due list rendered live from Customer Risk page. Reviews stop being a backlog.

### STR DRAFTING

Narrative auto-drafted from the alert evidence chain. Analyst reviews and signs — doesn't re-type. Filing-latency p95 is a live metric.

### BOI REFRESH

BOI Workflow page shows every reporting-company customer + freshness state. Stale BOI surfaces as an alert, not as someone's spreadsheet.

### BACKLOG STATUS

SLA timer + escalation engine. Per-queue breach-rate badges. Red queues become loud, not slide-decoration. CCO sees real status first.

**KEY TAKEAWAY** The unmeasured win isn't accuracy — it's the *analyst hours* the team gets back. Multiply across 100 analysts, 50 weeks: same dollars, signal instead of motion.

**Source:** industry false-positive baseline · Flagright + NICE Actimize 2026 · STR-filing-latency p95 cited in 2024 BSA consent orders · framework workflows in *Investigations*, *BOI Workflow*, *Tuning Lab*, *Audit & Evidence* dashboard pages

# The other half of the win is political. When the spec is the artifact, cross-team coordination collapses.

Most of an AML program's calendar is consumed by handoffs — 1LoD ⇄ 2LoD ⇄ IT ⇄ MRM ⇄ Audit. None of that produces detection. It produces alignment, and only when the alignment holds.

WHY THIS MATTERS

*"Examiner testing only determined whether controls existed rather than if they were actually being used."*

Across the largest 2024–25 enforcement actions, the same pattern recurs: **the controls existed; the coordination didn't.** Internal alerts get filed and never escalated. Rule changes wait on a calendar. Audit prep starts at the regulator's email. Removing the meetings removes the failure mode.

K&L GATES · 2024 BSA ENFORCEMENT-ACTION REVIEW · REFERENCED IN FCA DEAR-CEO MAR 2024

MONTHS → MINUTES

**Rule-change approval** → **one PR with cited regulation**

Today: months of 1LoD ⇄ 2LoD ⇄ IT ⇄ Audit alignment. With the framework: *git diff,* reviewer signs, audit chain records the change.

**BRA refresh** → **live render from spec**

Today: quarterly committee assembling spreadsheets. With the framework: the spec *is* the BRA — *Framework Alignment* dashboard shows it live.

**Annual MRM challenge** → **backtester self-serves**

Today: 4-week consultant engagement to validate scenarios. With the framework: per-rule precision/recall over 4 quarters, generated on demand.

**Pre-exam evidence assembly** → `aml audit-pack`

Today: 8-person sprint, 6 weeks, "what did we run?" forensics. With the framework: one CLI command exports the regulator-ready ZIP.

KEY TAKEAWAY Detection vendors sell better alerts. The framework removes the *meetings about the alerts* — and that is where the AML program's calendar actually lives.

**Source:** K&L Gates 2024 BSA enforcement-action review · FCA Dear-CEO Letter Mar 2024 · framework workflows in *Spec Editor*, *Framework Alignment*, *Tuning Lab*, *Audit Evidence* dashboard pages · `aml audit-pack` CLI

# Each pain has a single, evidenced answer — one file away.

Direct mapping from public-record regulator findings to specific framework capabilities. No black box, no vendor add-ons.

| | LEADER'S PAIN (REGULATOR-CITED) | | FRAMEWORK CAPABILITY (WHAT CHANGES) |
|---|---|---|---|
| 1 | *"We can't prove what we did."* | → | **Hash-chained** decisions.jsonl + deterministic replay. Any historical run reproducible byte-for-byte. |
| 2 | *"The backlog has been red on the board for years."* | → | **SLA timer + escalation engine.** Red queues become loud, not slide-decoration. Per-queue breach-rate badges visible to the board. |
| 3 | *"Our risk assessment is a PDF nobody trusts."* | → | **The spec IS the risk assessment** — machine-readable, cited line-by-line. *Framework Alignment* dashboard renders the live BRA from the spec. |
| 4 | *"My monitoring system is a model and I cannot validate it."* | → | **Per-rule MRM dossier + 4-quarter backtester.** "Is rule X still earning its keep?" answered before lunch — no vendor study. |

**Source:** Pain points 1, 2, 4, 6 from *docs/research/2026-04-aml-process-pain.md* · capability mapping in framework README + audit-evidence.md

# The framework arrives built, not blueprinted.

Open-source · runs in your perimeter · no per-seat licence · every rule cites a regulation · re-runs deterministically.

**8**

CURATED TYPOLOGIES

Every line cited to a regulator publication. Day-one detection.

**4**

ACTIVE JURISDICTIONS

US (FinCEN) · CA (FINTRAC) · UK (FCA) · EU (EBA). FATF Travel Rule.

**26**

DASHBOARD PAGES

Persona-aware. Sidebar filters by role — CCO sees board signal, analyst sees queue.

**9**

REFERENCE SPECS

Schedule-I bank · community bank · VASP · cyber-fraud · trade-based ML · …

**1,632**

TESTS · ALL GREEN

5 CI jobs per pull request. Lint, unit, API, e2e dashboard, docker.

**SHA-256**

AUDIT CHAIN

Hash-chained decision log. Tampering breaks attestation, by construction.

**MIT**

LICENCE

Apache 2.0 / MIT compatible. Forkable. Runs entirely in your perimeter.

**$0**

PER-SEAT FEES

No vendor licence renewals. Effectiveness pack quantifies what you spent.

**Source:** Live counts from *main* branch as of 2026-04-29 · github.com/tomqwu/aml_open_framework

THREE PATHS FROM HERE

# The next step takes five minutes.

READ · 15 MIN

## Show me the research.

Open the 10-pain-point doc with primary regulator citations plus a 2024 enforcement case study. Frame is unbranded.

TRY · 5 MIN

## Run it on my desk.

Clone, install, run a real audit pack against the canonical spec. No vendor demo cycle.

TALK · 30 MIN

## Walk me through it.

Bring your spec or pain — pull-request a typology, open a Discussion, or tag the maintainer.

OPEN SOURCE · MIT

github.com/tomqwu/aml_open_framework

**Source:** AML Open Framework · primary-source research grounded in FCA / FinCEN / OSFI / FINTRAC public record